

Project Part 1

November 30, 2024

0.0.1 Part1 VGG16 Quantization Aware Training

```
[1]: # Import libraries and load data (CIFAR 10)

# system library
import os
import time
import shutil

# NN library
import torch
import torch.nn as nn

# datasets library
import torchvision
import torchvision.transforms as transforms

# model library
from models import vgg_quant
from models import quant_layer

# data loading
batch_size = 100
num_workers = 2
normalize = transforms.Normalize(mean=[0.491, 0.482, 0.447], std=[0.247, 0.243, ↵
↵0.262])

train_data = torchvision.datasets.CIFAR10(
    root='data',
    train=True,
    download=True,
    transform=transforms.Compose([
        transforms.RandomCrop(32, padding=4),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        normalize,
```

```

    ]))

test_data = torchvision.datasets.CIFAR10(
    root='data',
    train=False,
    download=True,
    transform=transforms.Compose([
        transforms.ToTensor(),
        normalize,
    ]))

train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
    ↪shuffle=True, num_workers=num_workers)
test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size,
    ↪shuffle=False, num_workers=num_workers)

```

Files already downloaded and verified
Files already downloaded and verified

```

[2]: # Define functions for training, validation etc.
print_freq = 100

def train(train_loader, model, criterion, optimizer, epoch):
    batch_time = AverageMeter()    ## at the begining of each epoch, this should
    ↪be reset
    data_time = AverageMeter()
    losses = AverageMeter()
    top1 = AverageMeter()

    # switch to train mode
    model.train()
    end = time.time()

    for i, (x_train, y_train) in enumerate(train_loader):
        # record data loading time
        data_time.update(time.time() - end)

        # compute output and loss
        x_train = x_train.cuda()
        y_train = y_train.cuda()
        output = model(x_train)
        loss = criterion(output, y_train)

        # measure accuracy and record loss
        prec = accuracy(output, y_train)[0]

```

```

        losses.update(loss.item(), x_train.size(0))
        top1.update(prec.item(), x_train.size(0))

        # compute gradient and do SGD step
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # output epoch time and loss
        batch_time.update(time.time() - end)
        end = time.time()

    if i % print_freq == 0:
        print('Epoch: [{0}] [{1}/{2}]\t'
              'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
              'Data {data_time.val:.3f} ({data_time.avg:.3f})\t'
              'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
              'Prec {top1.val:.3f}% ({top1.avg:.3f}%)'.format(
                  epoch, i, len(train_loader), batch_time=batch_time,
                  data_time=data_time, loss=losses, top1=top1))

def validate(test_loader, model, criterion):
    batch_time = AverageMeter()
    losses = AverageMeter()
    top1 = AverageMeter()

    # switch to evaluate mode
    model.eval()
    end = time.time()

    with torch.no_grad():
        for i, (x_test, y_test) in enumerate(test_loader):
            # compute output
            x_test = x_test.cuda()
            y_test = y_test.cuda()
            output = model(x_test)
            loss = criterion(output, y_test)

            # measure accuracy and record loss
            prec = accuracy(output, y_test)[0]
            losses.update(loss.item(), x_test.size(0))
            top1.update(prec.item(), x_test.size(0))

            # measure elapsed time
            batch_time.update(time.time() - end)
            end = time.time()

```

```

        if i % print_freq == 0: # This line shows how frequently print out
        ↳ the status. e.g., i%5 => every 5 batch, prints out
            print('Test: [{0}/{1}]\t'
                  'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
                  'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
                  'Prec {top1.val:.3f}% ({top1.avg:.3f}%)'.format(
                    i, len(test_loader), batch_time=batch_time, loss=losses,
                    top1=top1))

    print(' * Prec {top1.avg:.3f}% '.format(top1=top1))
    return top1.avg

def accuracy(output, target, topk=(1,)):
    """Computes the precision@k for the specified values of k"""
    maxk = max(topk)
    batch_size = target.size(0)

    _, pred = output.topk(maxk, 1, True, True) # topk(k, dim=None,
    ↳ largest=True, sorted=True)
                                                # will output (max value, its
    ↳ index)
    pred = pred.t() # transpose
    correct = pred.eq(target.view(1, -1).expand_as(pred)) # "-1": calculate
    ↳ automatically

    res = []
    for k in topk:
        correct_k = correct[:k].view(-1).float().sum(0) # view(-1): make a
    ↳ flattened 1D tensor
        res.append(correct_k.mul_(100.0 / batch_size)) # correct: size of
    ↳ [maxk, batch_size]
    return res

class AverageMeter(object):
    """Computes and stores the average and current value"""
    def __init__(self):
        self.reset()

    def reset(self):
        self.val = 0
        self.avg = 0
        self.sum = 0
        self.count = 0

```

```

def update(self, val, n=1):
    self.val = val
    self.sum += val * n    ## n is impact factor
    self.count += n
    self.avg = self.sum / self.count

def save_checkpoint(state, is_best, fdir):
    filepath = os.path.join(fdir, 'checkpoint.pth')
    torch.save(state, filepath)
    if is_best:
        shutil.copyfile(filepath, os.path.join(fdir, 'model_best.pth.tar'))

def adjust_learning_rate(optimizer, epoch):
    """For resnet, the lr starts from 0.1, and is divided by 10 at 80 and 120_
    epochs"""
    adjust_list = [150, 225]
    if epoch in adjust_list:
        for param_group in optimizer.param_groups:
            param_group['lr'] = param_group['lr'] * 0.1

```

```

[3]: # Configure model
model_name = 'project'
model_project = vgg_quant.VGG16_quant()

# Adjust certain layers
model_project.features[24] = quant_layer.QuantConv2d(256, 8, ▯
    ▯kernel_size=3,padding=1)
model_project.features[25] = nn.BatchNorm2d(8)
model_project.features[27] = quant_layer.QuantConv2d(8, 8, kernel_size=3, ▯
    ▯padding=1)
model_project.features[30] = quant_layer.QuantConv2d(8, 512, kernel_size=3, ▯
    ▯padding=1)
del model_project.features[28]

# parameters for training
lr = 0.02
weight_decay = 1e-4
epochs = 100
best_prec = 0

model_project = model_project.cuda()
criterion = nn.CrossEntropyLoss().cuda()
optimizer = torch.optim.SGD(model_project.parameters(), lr=lr, momentum=0.8, ▯
    ▯weight_decay=weight_decay)

```

```

# saving path
if not os.path.exists('result'):
    os.makedirs('result')
fdir = 'result/'+str(model_name)
if not os.path.exists(fdir):
    os.makedirs(fdir)

```

```

[ ]: # Train and validate 4bit vgg16 model
for epoch in range(100, 100+epochs):
    adjust_learning_rate(optimizer, epoch)
    train(train_loader, model_project, criterion, optimizer, epoch)

    # evaluate on test set
    print("Validation starts")
    prec = validate(test_loader, model_project, criterion)

    # remember best precision and save checkpoint
    is_best = prec > best_prec
    best_prec = max(prec, best_prec)
    print('best acc: {:.1f}'.format(best_prec))
    save_checkpoint({
        'epoch': epoch + 1,
        'state_dict': model_project.state_dict(),
        'best_prec': best_prec,
        'optimizer': optimizer.state_dict(),
    }, is_best, fdir)

```

```

[4]: # Validate 4bit vgg16 model on test dataset
fdir = 'result/'+str(model_name)+'/' + 'model_best.pth.tar'
checkpoint = torch.load(fdir)
model_project.load_state_dict(checkpoint['state_dict'])

criterion = nn.CrossEntropyLoss().cuda()
model_project.eval()
model_project.cuda()
prec = validate(test_loader, model_project, criterion)

```

```

Test: [0/100]   Time 1.075 (1.075)      Loss 0.3422 (0.3422)      Prec 92.000%
(92.000%)
* Prec 92.140%

```

```

[5]: # Prehook
class SaveOutput:
    def __init__(self):
        self.outputs = []
    def __call__(self, module, module_in):

```

```

        self.outputs.append(module_in)
    def clear(self):
        self.outputs = []

save_output = SaveOutput()
for layer in model_project.modules():
    if isinstance(layer, torch.nn.Conv2d):
        # print("prehooked")
        layer.register_forward_pre_hook(save_output)

dataiter = iter(train_loader)
images, labels = next(dataiter)
images = images.cuda()
out = model_project(images)

print("feature 27th layer's input size:", save_output.outputs[8][0].size())
print("feature 29th layer's input size:", save_output.outputs[9][0].size())

```

feature 27th layer's input size: torch.Size([100, 8, 4, 4])
feature 29th layer's input size: torch.Size([100, 8, 4, 4])

```

[6]: # Find x_int and w_int for the 8*8 convolution layer
layer = model_project.features[27]
x = save_output.outputs[8][0]

w_bits = 4
w_alpha = layer.weight_quant.wgt_alpha
w_delta = w_alpha/(2**(w_bits-1)-1)
weight_q = layer.weight_q # quantized value is stored during the training
weight_int = weight_q/w_delta

x_bits = 4
x_alpha = layer.act_alpha
x_delta = x_alpha/(2**(x_bits-1)-1)
act_quant_fn = quant_layer.act_quantization(x_bits) # define the quantization
    ↪function
x_q = act_quant_fn(x, x_alpha) # create the quantized value for x
x_int = x_q/x_delta

```

```

[7]: # Check the recovered p_sum has similar value to the un-quantized original
    ↪output
conv_int = torch.nn.Conv2d(in_channels=8, out_channels=8, kernel_size=3,
    ↪padding=1, bias=False)
conv_int.weight = torch.nn.parameter.Parameter(weight_int)
output_int = conv_int(x_int)
output_int = model_project.features[28](output_int)
output_recovered = output_int * w_delta * x_delta

```

```
# calculate the differences  
print((output_recovered - save_output.outputs[9][0]).sum())
```

```
tensor(0.0006, device='cuda:0', grad_fn=<SumBackward0>)
```