Project Part Alpha Pruning

December 1, 2024

0.0.1 Project Part Alpha Pruning

[1]:
```
# Import libraries and load data (CIFAR 10)

# system library
import os
import time
import shutil

# NN library
import torch
import torch.nn as nn
import torch.nn.utils.prune as prune

# datasets library
import torchvision
import torchvision.transforms as transforms

# model library
from models import vgg_quant
from models import quant_layer


# data loading
batch_size = 100
num_workers = 2
normalize = transforms.Normalize(mean=[0.491, 0.482, 0.447], std=[0.247, 0.243,␣
 ↪0.262])


train_data = torchvision.datasets.CIFAR10(
    root='data',
    train=True,
    download=True,
    transform=transforms.Compose([
        transforms.RandomCrop(32, padding=4),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
```

```
        normalize,
    ]))


test_data = torchvision.datasets.CIFAR10(
    root='data',
    train=False,
    download=True,
    transform=transforms.Compose([
        transforms.ToTensor(),
        normalize,
    ]))


train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,␣
 ↪shuffle=True, num_workers=num_workers)
test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size,␣
 ↪shuffle=False, num_workers=num_workers)
```

```
Files already downloaded and verified
Files already downloaded and verified
```

[2]:
```
# Define functions for training, validation etc.
print_freq = 100

def train(train_loader, model, criterion, optimizer, epoch):
    batch_time = AverageMeter()   ## at the begining of each epoch, this should␣
 ↪be reset
    data_time = AverageMeter()
    losses = AverageMeter()
    top1 = AverageMeter()

    # switch to train mode
    model.train()
    end = time.time()

    for i, (x_train, y_train) in enumerate(train_loader):
        # record data loading time
        data_time.update(time.time() - end)

        # compute output and loss
        x_train = x_train.cuda()
        y_train = y_train.cuda()
        output = model(x_train)
        loss = criterion(output, y_train)

        # measure accuracy and record loss
```

```
        prec = accuracy(output, y_train)[0]
        losses.update(loss.item(), x_train.size(0))
        top1.update(prec.item(), x_train.size(0))

        # compute gradient and do SGD step
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # output epoch time and loss
        batch_time.update(time.time() - end)
        end = time.time()

        if i % print_freq == 0:
            print('Epoch: [{0}][{1}/{2}]\t'
                  'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
                  'Data {data_time.val:.3f} ({data_time.avg:.3f})\t'
                  'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
                  'Prec {top1.val:.3f}% ({top1.avg:.3f}%)'.format(
                   epoch, i, len(train_loader), batch_time=batch_time,
                   data_time=data_time, loss=losses, top1=top1))


def validate(test_loader, model, criterion):
    batch_time = AverageMeter()
    losses = AverageMeter()
    top1 = AverageMeter()

    # switch to evaluate mode
    model.eval()
    end = time.time()

    with torch.no_grad():
        for i, (x_test, y_test) in enumerate(test_loader):
            # compute output
            x_test = x_test.cuda()
            y_test = y_test.cuda()
            output = model(x_test)
            loss = criterion(output, y_test)

            # measure accuracy and record loss
            prec = accuracy(output, y_test)[0]
            losses.update(loss.item(), x_test.size(0))
            top1.update(prec.item(), x_test.size(0))

            # measure elapsed time
            batch_time.update(time.time() - end)
```

```
            end = time.time()

            if i % print_freq == 0:  # This line shows how frequently print out␣
 ↪the status. e.g., i%5 => every 5 batch, prints out
                print('Test: [{0}/{1}]\t'
                  'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
                  'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
                  'Prec {top1.val:.3f}% ({top1.avg:.3f}%)'.format(
                   i, len(test_loader), batch_time=batch_time, loss=losses,
                   top1=top1))

    print(' * Prec {top1.avg:.3f}% '.format(top1=top1))
    return top1.avg


def accuracy(output, target, topk=(1,)):
    """Computes the precision@k for the specified values of k"""
    maxk = max(topk)
    batch_size = target.size(0)

    _, pred = output.topk(maxk, 1, True, True) # topk(k, dim=None,␣
 ↪largest=True, sorted=True)
                                               # will output (max value, its␣
 ↪index)
    pred = pred.t()                                          # transpose
    correct = pred.eq(target.view(1, -1).expand_as(pred))    # "-1": calculate␣
 ↪automatically

    res = []
    for k in topk:
        correct_k = correct[:k].view(-1).float().sum(0)  # view(-1): make a␣
 ↪flattened 1D tensor
        res.append(correct_k.mul_(100.0 / batch_size))   # correct: size of␣
 ↪[maxk, batch_size]
    return res


class AverageMeter(object):
    """Computes and stores the average and current value"""
    def __init__(self):
        self.reset()

    def reset(self):
        self.val = 0
        self.avg = 0
        self.sum = 0
```

```
        self.count = 0

    def update(self, val, n=1):
        self.val = val
        self.sum += val * n    ## n is impact factor
        self.count += n
        self.avg = self.sum / self.count


def save_checkpoint(state, is_best, fdir):
    filepath = os.path.join(fdir, 'checkpoint.pth')
    torch.save(state, filepath)
    if is_best:
        shutil.copyfile(filepath, os.path.join(fdir, 'model_best.pth.tar'))


def adjust_learning_rate(optimizer, epoch):
    """For resnet, the lr starts from 0.1, and is divided by 10 at 80 and 120␣
 ↪epochs"""
    adjust_list = [150, 225]
    if epoch in adjust_list:
        for param_group in optimizer.param_groups:
            param_group['lr'] = param_group['lr'] * 0.1
```

[3]:
```
# Configure model
model_name = 'project'
model_project = vgg_quant.VGG16_quant()

# Adjust certain layers
model_project.features[24] = quant_layer.QuantConv2d(256, 8,␣
 ↪kernel_size=3,padding=1)
model_project.features[25] = nn.BatchNorm2d(8)
model_project.features[27] = quant_layer.QuantConv2d(8, 8, kernel_size=3,␣
 ↪padding=1)
model_project.features[30] = quant_layer.QuantConv2d(8, 512, kernel_size=3,␣
 ↪padding=1)
del model_project.features[28]

# parameters for training
lr = 0.02
weight_decay = 1e-4
epochs = 100
best_prec = 0

model_project = model_project.cuda()
criterion = nn.CrossEntropyLoss().cuda()
```

```
optimizer = torch.optim.SGD(model_project.parameters(), lr=lr, momentum=0.8,␣
 ↪weight_decay=weight_decay)

# saving path
if not os.path.exists('result'):
    os.makedirs('result')
fdir = 'result/'+str(model_name)
if not os.path.exists(fdir):
    os.makedirs(fdir)
```

[4]:
```
# Validate 4bit vgg16 model on test dataset
fdir = 'result/'+str(model_name)+'/model_best.pth.tar'
checkpoint = torch.load(fdir)
model_project.load_state_dict(checkpoint['state_dict'])

criterion = nn.CrossEntropyLoss().cuda()
model_project.eval()
model_project.cuda()
prec = validate(test_loader, model_project, criterion)
```

```
Test: [0/100]   Time 1.092 (1.092)      Loss 0.3422 (0.3422)    Prec 92.000%
(92.000%)
 * Prec 92.140%
```

Unstructured pruning with 80% sparsity for the modified conv layer

[5]:
```
# Unstructured pruning
# Structured pruning
layer = model_project.features[27]
prune.ln_structured(layer, name='weight', amount=0.8, dim=1, n=1)
```

[5]:
```
QuantConv2d(
  8, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
  (weight_quant): weight_quantize_fn()
)
```

[6]:
```
# Check sparsity
mask1 = model_project.features[27].weight_mask
sparsity_mask1 = (mask1 == 0).sum() / mask1.nelement()
print("Sparsity level: ", sparsity_mask1)

# Check precision before training
prec = validate(test_loader, model_project, criterion)
```

```
Sparsity level:  tensor(0.7500, device='cuda:0')
Test: [0/100]   Time 0.441 (0.441)      Loss 4.5215 (4.5215)    Prec 26.000%
(26.000%)
 * Prec 26.300%
```

[ ]:
```
# Retrain
for epoch in range(10):
    adjust_learning_rate(optimizer, epoch)
    train(train_loader, model_project, criterion, optimizer, epoch)

    # evaluate on test set
    print("Validation starts")
    prec = validate(test_loader, model_project, criterion)

    # remember best precision and save checkpoint
    is_best = prec > best_prec
    best_prec = max(prec,best_prec)
```

[8]:
```
# Check precision after training
prec = validate(test_loader, model_project, criterion)
```

```
Test: [0/100]   Time 0.356 (0.356)      Loss 0.3956 (0.3956)    Prec 86.000%
(86.000%)
 * Prec 89.700%
```

Structured pruning with 40% sparsity for the modified conv layer

[9]:
```
# Reload checkpoint
model_project = vgg_quant.VGG16_quant()
model_project.features[24] = quant_layer.QuantConv2d(256, 8,␣
 ↪kernel_size=3,padding=1)
model_project.features[25] = nn.BatchNorm2d(8)
model_project.features[27] = quant_layer.QuantConv2d(8, 8, kernel_size=3,␣
 ↪padding=1)
model_project.features[30] = quant_layer.QuantConv2d(8, 512, kernel_size=3,␣
 ↪padding=1)
del model_project.features[28]

fdir = 'result/'+'project'+'/model_best.pth.tar'
checkpoint = torch.load(fdir)
model_project.load_state_dict(checkpoint['state_dict'])
model_project.cuda()

# Structured pruning
layer = model_project.features[27]
prune.ln_structured(layer, name='weight', amount=0.4, dim=1, n=1)
```

[9]:
```
QuantConv2d(
  8, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
  (weight_quant): weight_quantize_fn()
)
```

[10]:
```
# Check sparsity
mask1 = model_project.features[27].weight_mask
sparsity_mask1 = (mask1 == 0).sum() / mask1.nelement()
print("Sparsity level: ", sparsity_mask1)

# Check precision before training
prec = validate(test_loader, model_project, criterion)
```

```
Sparsity level:  tensor(0.3750, device='cuda:0')
Test: [0/100]   Time 0.374 (0.374)      Loss 1.5509 (1.5509)    Prec 66.000%
(66.000%)
 * Prec 57.910%
```

[ ]:
```
# Retrain
for epoch in range(10):
    adjust_learning_rate(optimizer, epoch)
    train(train_loader, model_project, criterion, optimizer, epoch)

    # evaluate on test set
    print("Validation starts")
    prec = validate(test_loader, model_project, criterion)

    # remember best precision and save checkpoint
    is_best = prec > best_prec
    best_prec = max(prec,best_prec)
```

[12]:
```
# Check precision after training
prec = validate(test_loader, model_project, criterion)
```

```
Test: [0/100]   Time 0.331 (0.331)      Loss 0.8657 (0.8657)    Prec 83.000%
(83.000%)
 * Prec 73.840%
```